## Measured MCTS generation work at the frozen source boundary

All shard work, including pending targets; unequal work. No matched-compute gain claimed.

| Quantity | Recorded value |
|---|---|
| Newly generated output tokens | 852,951,580 |
| Cached output tokens (excluded from generation charge) | 184,560,843 |
| Unique scheduled model calls | 8,035,879 |
| CPU executor thread seconds (additive across threads) | 46,502.3 |
| Distinct A/B-supported complete reactions | 16 |
| Targets with every essential route reaction A/B | 1 |
| Supported reactions per million generated tokens | 0.0187584 |
| Covered targets per million generated tokens | 0.0011724 |
| GPU-active seconds | Unknown |

No conversion of cache-return tokens or summed worker elapsed time into fictitious GPU time.

Source-checking coverage is incomplete; these efficiency values are evidence-discovery lower bounds.