Measured MCTS generation work at the frozen source boundary

All shard work, including pending targets; unequal work. No matched-compute gain claimed.

Quantity	Recorded value
Newly generated output tokens	852,951,580
Cached output tokens (excluded from generation charge)	184,560,843
Unique scheduled model calls	8,035,879
CPU executor thread seconds (additive across threads)	46,502.3
Distinct A/B-supported complete reactions	27
Targets with every essential route reaction A/B	1
Supported reactions per million generated tokens	0.0316548
Covered targets per million generated tokens	0.0011724
GPU-active seconds	Unknown

No conversion of cache-return tokens or summed worker elapsed time into fictitious GPU time.

Source-checking coverage is incomplete; these efficiency values are evidence-discovery lower bounds.